

1. Download all the data in this folder <https://drive.google.com/open?id=1Z4TyI7FcFVEx8qdl...>
path/to/the/image.tif, category

where the categories are numbered 0 to 15, in the following order:

- 0 letter
- 1 form
- 2 email
- 3 handwritten
- 4 advertisement
- 5 scientific report
- 6 scientific publication
- 7 specification
- 8 file folder
- 9 news article
- 10 budget
- 11 invoice
- 12 presentation
- 13 questionnaire
- 14 resume
- 15 memo

2. On this image data, you have to train 3 types of models as given below. You have to spli

3. Try not to load all the images into memory, use the gernerators that we have given the
or you can use this method also

<https://medium.com/@vijayabhaskar96/tutorial-on-keras-imagedatagenerator-with-flow-from-da>


<https://medium.com/@vijayabhaskar96/tutorial-on-keras-flow-from-dataframe-1fd4493d237c>

4. You are free to choose Learning rate, optimizer, loss function, image augmentation, any

5. Use tensorboard for every model and analyse your gradients. (you need to upload the scr

Note: fit_genarator() method will have problems with the tensorboard histograms, try to del

6. You can check about Transfer Learning in this link - <https://blog.keras.io/building-pow>



```
#!wget --header="Host: doc-08-6g-docs.googleusercontent.com" --header="User-Agent: Mozilla/5.
```

```

get_ipython().system_raw("unrar x rv1-cdip.rar")

%load_ext tensorboard

from keras_preprocessing.image import ImageDataGenerator
from tensorflow import keras
from sklearn.model_selection import train_test_split
import pandas as pd

df = pd.read_csv('/content/labels_final.csv')

train_data, valid_data = train_test_split(df, test_size=0.3, random_state=42)

#data augmentation
datagen = ImageDataGenerator(rescale=1./255)

#get images from different directories
train_generator = datagen.flow_from_dataframe(dataframe = train_data,
                                              directory="data_final", x_col="path", y_col="label",
                                              class_mode="raw", target_size=(256,256), batch_size=32)

valid_generator = datagen.flow_from_dataframe(dataframe = valid_data,
                                              directory="data_final", x_col="path", y_col="label",
                                              class_mode="raw", target_size=(256,256), batch_size=32)

📄 Found 33600 validated image filenames.
   Found 14400 validated image filenames.

```

▼ Model-1

1. Use [VGG-16](#) pretrained network without Fully Connected layers and initialize all the weights
2. After VGG-16 network without FC layers, add a new Conv block (1 Conv layer and 1 Maxpooling layer)
3. Final architecture will be **INPUT --> VGG-16 without Top layers(FC) --> Conv Layer --> Maxpooling layer --> FC layers --> Output layer**
4. Train only new Conv block, FC layers, output layer. Don't train the VGG-16 network.

```

#Model 1
import datetime
from keras.models import Model

vgg_model = keras.applications.VGG16(weights='imagenet',
                                     include_top=False,

```

```
input_shape=(256, 256, 3))
```

```
x = vgg_model.output
x = keras.layers.Conv2D(32, (3, 3),padding = 'same')(x)
x = keras.layers.MaxPooling2D(pool_size=(2, 2))(x)
x = keras.layers.Flatten()(x)
x = keras.layers.Dense(128, activation='relu')(x)
x = keras.layers.Dropout(.5)(x)
x = keras.layers.Dense(64, activation='relu')(x)
x = keras.layers.Dropout(.5)(x)
preds = keras.layers.Dense(16, activation='softmax')(x)
model1 = Model(inputs = vgg_model.input,outputs=preds)
```

☞ Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/vgg1>
58892288/58889256 [=====] - 0s 0us/step

```
for layer in model1.layers[:19]:
    layer.trainable=False
for layer in model1.layers[19:]:
    layer.trainable=True
```

```
model1.summary()
```

☞

Model: "functional_1"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0

```
import tensorflow as tf
model1.compile(optimizer= 'adam',loss='sparse_categorical_crossentropy',metrics=['accuracy'])

#tensorflow callback
log_dir="logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = keras.callbacks.TensorBoard(log_dir=log_dir,histogram_freq=1, write_graphs=True)

#early stopping callback
early_stopping = keras.callbacks.EarlyStopping(monitor='val_accuracy', patience=2)

STEP_SIZE_TRAIN=train_generator.n//train_generator.batch_size
STEP_SIZE_VALID=valid_generator.n//valid_generator.batch_size

with tf.device('/device:GPU:0'):
    model1.fit_generator(generator=train_generator,
                        steps_per_epoch=STEP_SIZE_TRAIN,
                        validation_data=valid_generator,
                        validation_steps=STEP_SIZE_VALID,
                        callbacks = [tensorboard_callback, early_stopping],
                        epochs=5)
```

```

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `TensorBoard
WARNING:tensorflow:From <ipython-input-10-f5fb2a2567c6>:21: Model.fit_generator (from t
Instructions for updating:
Please use Model.fit, which supports generators.
Epoch 1/5
   1/1050 [.....] - ETA: 0s - loss: 3.3809 - accuracy: 0.0312W
Instructions for updating:
use `tf.profiler.experimental.stop` instead.
1050/1050 [=====] - 541s 516ms/step - loss: 2.0371 - accuracy:
Epoch 2/5
1050/1050 [=====] - 541s 515ms/step - loss: 1.5413 - accuracy:
Epoch 3/5
1050/1050 [=====] - 540s 515ms/step - loss: 1.3827 - accuracy:
Epoch 4/5
1050/1050 [=====] - 540s 515ms/step - loss: 1.2765 - accuracy:
Epoch 5/5
1050/1050 [=====] - 541s 515ms/step - loss: 1.1900 - accuracy:

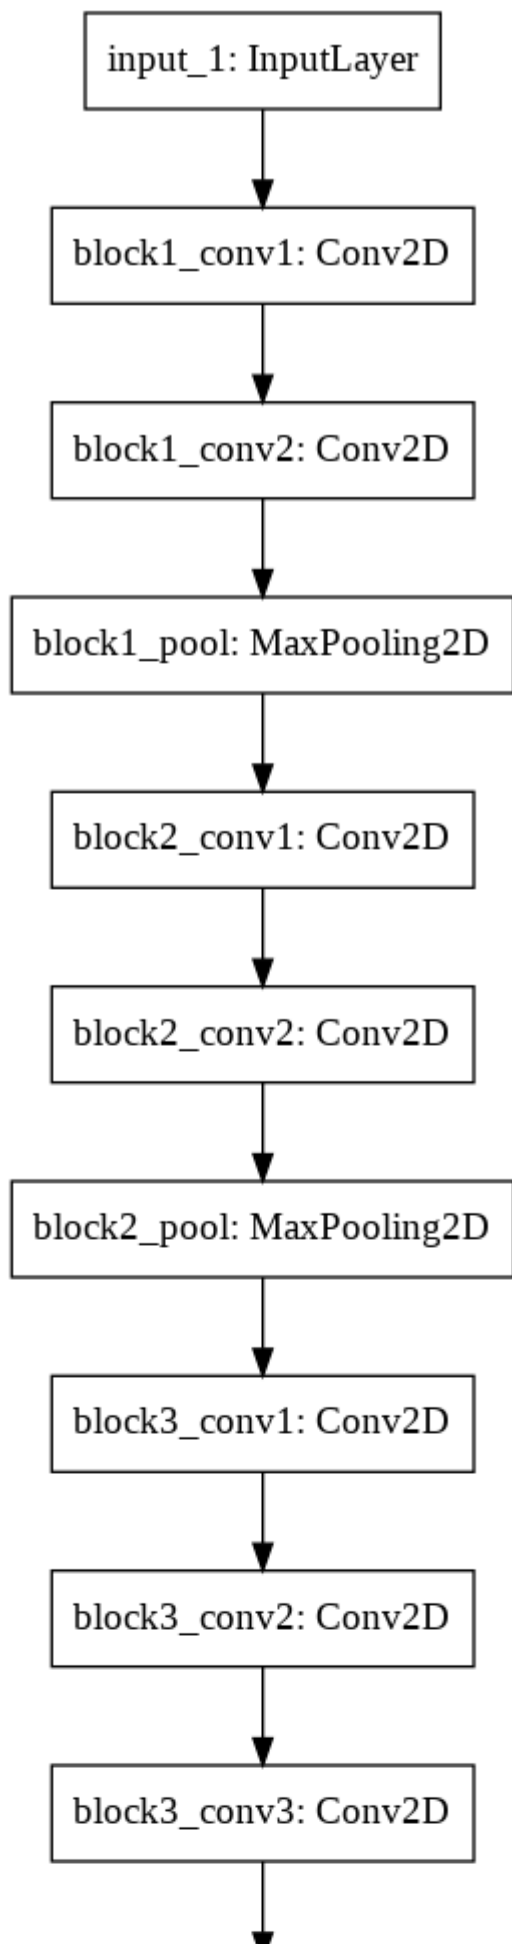
```

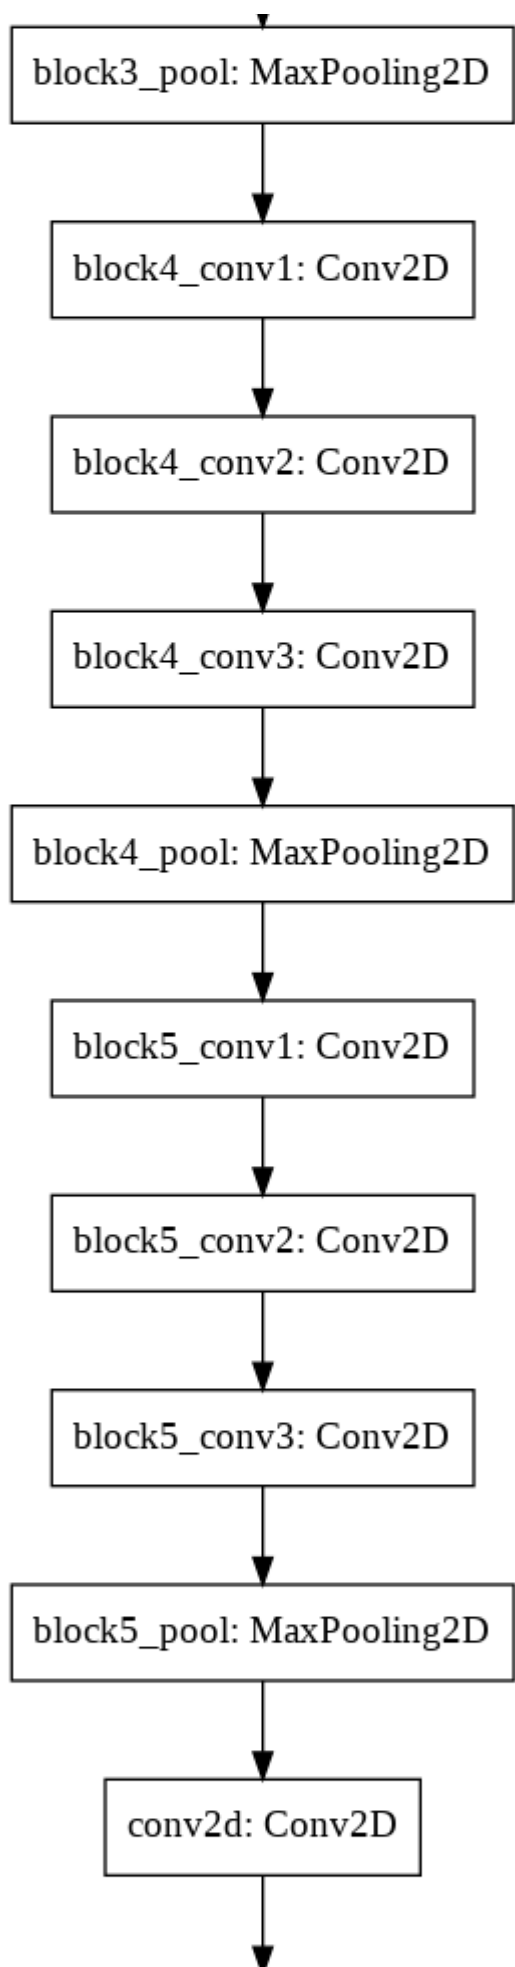
```


keras.utils.plot_model(
    model1, to_file='model1.png', show_shapes=False, show_layer_names=True,
    rankdir='TB', expand_nested=False, dpi=96
)

```









```
#!/rm -rf ./logs
```

```
|
```

```
%tensorboard --logdir /content/logs/fit/20200827-164553
```




```
for layer in model_2.layers[:13]:  
    layer.trainable=False  
for layer in model_2.layers[13:]:  
    layer.trainable=True  
  
model_2.summary()
```



- The dataset difference between the pretrained model and our models are also high regarding sizes, variance of categories.

TensorBoard


SCALARS

GRAPHS

INACTIVE

▼ Observations:


- The validation accuracy for Model 1 = 0.69
- The imagenet dataset contains 14 million images and our model trained on 33600 images, a huge difference in dataset
- Maybe the imagenet dataset has less category of file type data



▼ Model-2



1. Use [VGG-16](#) pretrained network without Fully Connected layers and initialize all the weights
2. After VGG-16 network without FC layers, don't use FC layers, use conv layers only as Fully Connected layers
3. Final architecture will be VGG-16 without FC layers(without top), 2 Conv layers identical to FC layers, 1 output layer. Don't train the VGG-16 network



```
vgg_model = keras.applications.VGG16(weights='imagenet',
                                     include_top=False,
                                     input_shape=(256, 256, 3))
```

```
x2 = vgg_model.output
x2 = keras.layers.Conv2D(4096, (8, 8), padding='valid')(x2)
x2 = keras.layers.Conv2D(4096, (1, 1), padding='valid')(x2)
x2 = keras.layers.Flatten()(x2)
preds = keras.layers.Dense(16, activation='softmax')(x2)
```

```
model_2 = Model(inputs = vgg_model.input, outputs=preds)
```

```
for layer in model_2.layers[:19]:
    layer.trainable=False
```

```
for layer in model_2.layers[19:]:
    layer.trainable=True
```

```
model_2.summary()
```



Model: "functional_3"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0
conv2d_1 (Conv2D)	(None, 1, 1, 4096)	134221824
conv2d_2 (Conv2D)	(None, 1, 1, 4096)	16781312
flatten_1 (Flatten)	(None, 4096)	0

```
model_2.compile(optimizer= 'adam',loss='sparse_categorical_crossentropy',metrics=['accuracy'])
```

```
#tensorflow callback
```

```
log_dir="logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
```

```
tensorboard_callback = keras.callbacks.TensorBoard(log_dir=log_dir,histogram_freq=1, write_gr
```

```
#early stopping callback
```

```
early_stopping = keras.callbacks.EarlyStopping(monitor='val_accuracy', patience=2)
```

```
early_stopping_monitor=EarlyStopping(monitor='val_accuracy', patience=2,
```

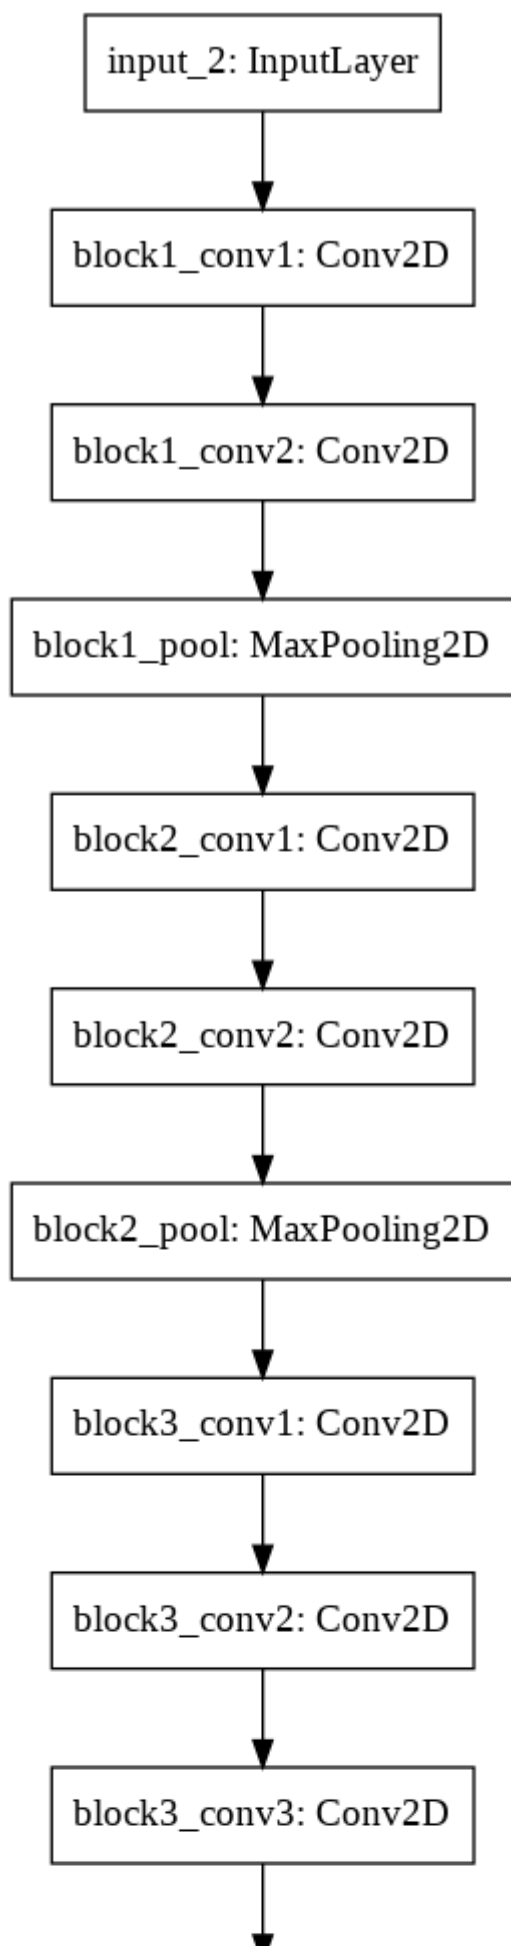
```
model_2.fit_generator(generator=train_generator,
                      steps_per_epoch=STEP_SIZE_TRAIN,
                      validation_data=valid_generator,
                      validation_steps=STEP_SIZE_VALID,
                      callbacks = [tensorboard_callback, early_stopping],
                      epochs=5)
```

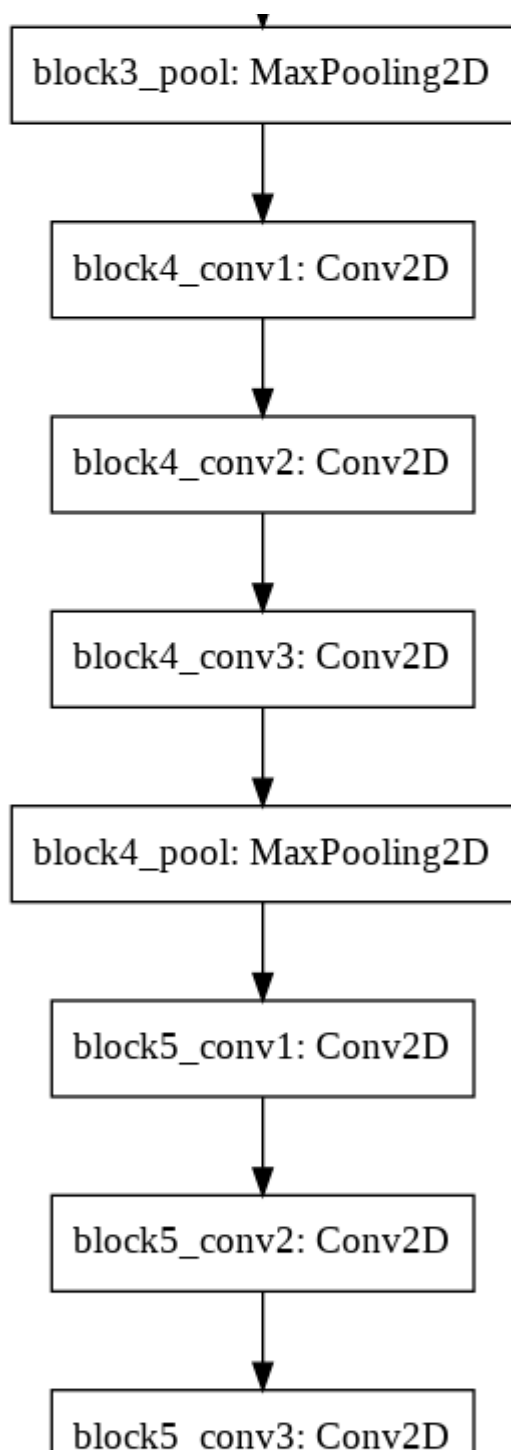
↳ `flow: `write_grads` will be ignored in TensorFlow 2.0 for the `TensorBoard` Callback.`

```
=====] - 1096s 1s/step - loss: 15.2254 - accuracy: 0.5221 - val_lo
=====] - 1091s 1s/step - loss: 1.7481 - accuracy: 0.6596 - val_lo
=====] - 1093s 1s/step - loss: 0.9951 - accuracy: 0.7239 - val_lo
=====] - 1094s 1s/step - loss: 125.0922 - accuracy: 0.6199 - val_l
thon.keras.callbacks.History at 0x7f17ceb93588>
```

```
keras.utils.plot_model(
    model_2, to_file='model2.png', show_shapes=False, show_layer_names=True,
    rankdir='TB', expand_nested=False, dpi=96)
```

↳





```
%tensorboard --logdir logs/fit/20200827-173114
```



TensorBoard

SCALARS

GRAPHS

DIS

INACTIVE

- ☐ Show data download links
- ☐ Ignore outliers in chart scaling

Tooltip sorting method: default

Smoothing



0.6

Horizontal Axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs


☐ ☐ train

☐ ☐ validation

TOGGLE ALL RUNS


epoch_accuracy

epoch_accuracy



epoch_loss

epoch_loss



Observations:

- In Model 2 the validation accuracy = 0.64
- Even though the only difference in vgg_16 and this model was last two layers in vgg-16 were dense and model 2 had convoluted data
- Maybe imagenet data doesn't consist variety of file type images

Model-3

1. Use same network as Model-2 'INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers :

Model: "functional_3"

Layer (type)	Output Shape	Param #
=====		
input 2 (InputLayer)	[(None, 256, 256, 3)]	0

```
model_2.compile(optimizer= 'adam',loss='sparse_categorical_crossentropy',metrics=['accuracy'])
```

```
#tensorflow callback
```

```
log_dir="logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
```

```
tensorboard_callback = keras.callbacks.TensorBoard(log_dir=log_dir,histogram_freq=1, write_g
```

```
#early stopping callback
```

```
early_stopping = keras.callbacks.EarlyStopping(monitor='val_accuracy', patience=2)
```

```
model_2.fit_generator(generator=train_generator,
                      steps_per_epoch=STEP_SIZE_TRAIN,
                      validation_data=valid_generator,
                      validation_steps=STEP_SIZE_VALID,
                      callbacks = [tensorboard_callback, early_stopping],
                      epochs=5)
```

☞ tensorflow: `write_grads` will be ignored in TensorFlow 2.0 for the `TensorBoard` Callback.

```
=====] - 1652s 2s/step - loss: 4.9436 - accuracy: 0.0635 - val_
```

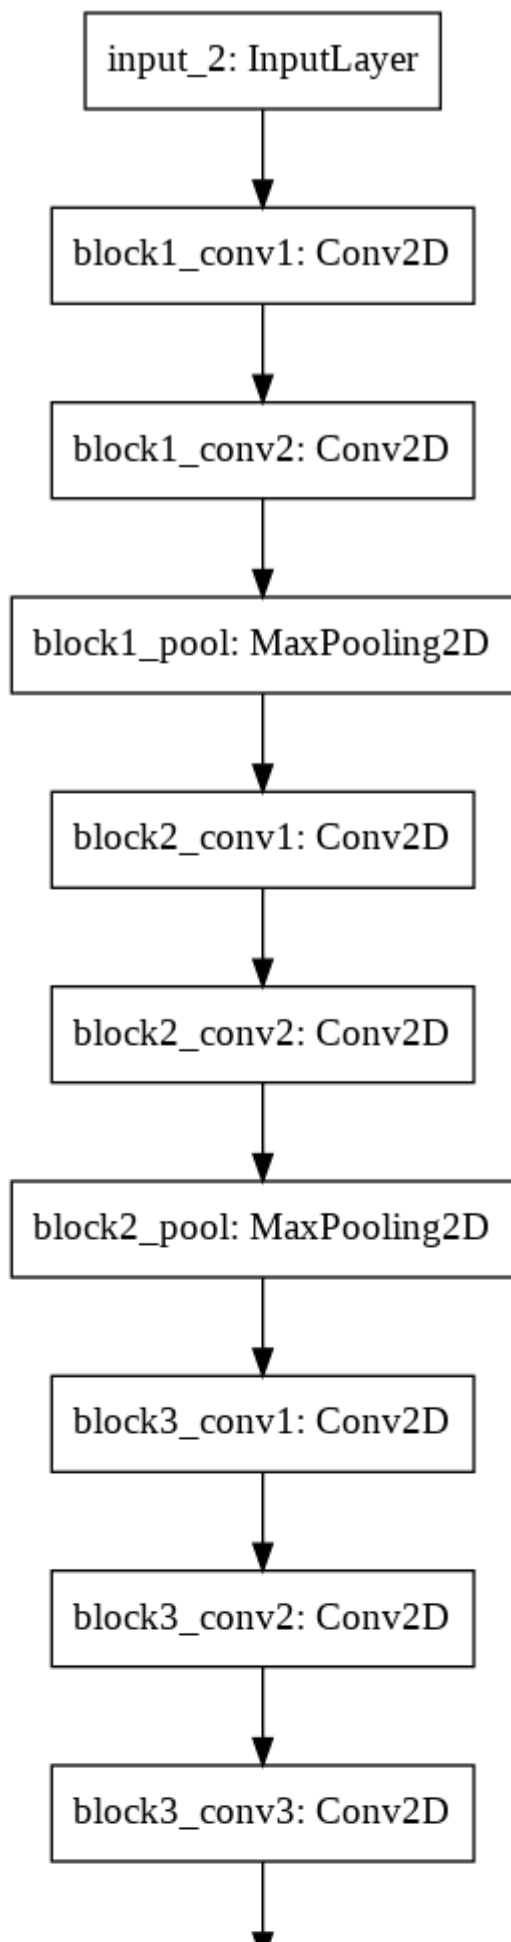
```
=====] - 1643s 2s/step - loss: 2.7900 - accuracy: 0.0634 - val_
```

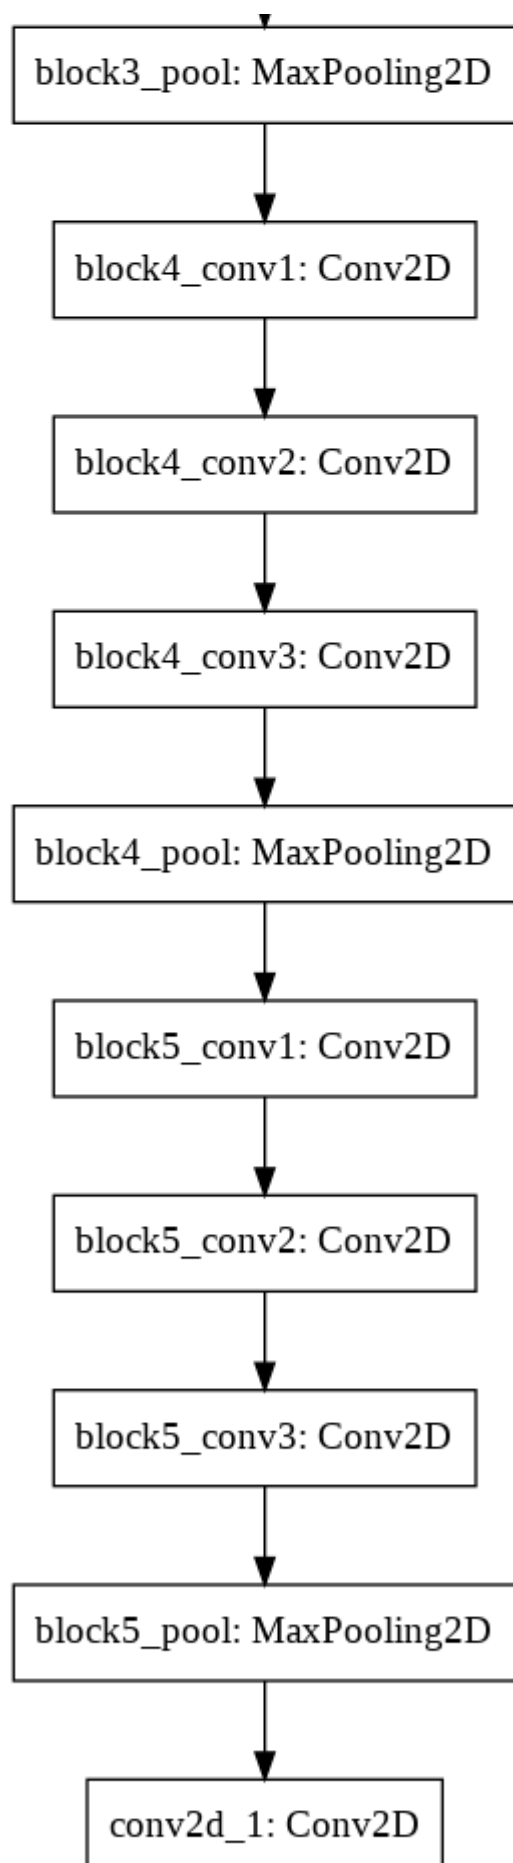
```
=====] - 1652s 2s/step - loss: 2.7869 - accuracy: 0.0625 - val_
```

```
python.keras.callbacks.History at 0x7f17cc1229b0>
```

```
keras.utils.plot_model(
    model_2, to_file='model_2.png', show_shapes=False, show_layer_names=True,
    rankdir='TB', expand_nested=False, dpi=96)
```

☞





%tensorboard --logdir /content/logs/fit/20200827-184416



TensorBoard

SCALARS

GRAPHS

DIS

INACTIVE

☐ Show data download links☐ Ignore outliers in chart scalingTooltip sorting
method: default

Smoothing

0.6

Horizontal Axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs

☐ train☐ validation


TOGGLE ALL RUNS

/content/logs/fit/20200827-184416

Filter tags (regular expressions supported)


epoch_accuracy

epoch_accuracy



epoch_loss

epoch_loss



Observations:

- The validation accuracy for model 3 is 0.0619
- In these model we have also trained the pre trained layers of vgg-16 models
- The accuracy is so low than other two models maybe the weights generated by pretrained models were not useful for the images we have.